

Reconnaissance Techniques and Subdomain Enumeration

Tiago Almeida

October 1, 2025

Abstract

This report presents the resolution of the exercises from the second class of the *Theory and Practice of Security Attacks* course. The assignment consisted of applying reconnaissance techniques and tools to the domain **chime.com**. The following sections document the methodology, commands, results, and findings from passive and active subdomain enumeration, HTTP service discovery, URL scanning, and Google dorking performed against the target domain.

Contents

1	Passive subdomain enumeration	2
1.1	Tools used	2
1.2	Methodology	2
1.3	Commands	2
1.4	Results	3
2	Active subdomain enumeration	3
2.1	Tools and wordlists	3
2.2	Commands	4
2.3	Results and analysis	4
3	HTTP(S) service discovery	5
3.1	Tools	5
3.2	Commands (examples)	5
4	URL scanning (detailed analysis)	5
4.1	Top 5 Identified Technologies	6
4.2	Content discovery	7

1 Passive subdomain enumeration

1.1 Tools used

For this exercise I used both recommended tools: **subfinder** and **assetfinder**. The reasons for using both were:

- to learn how to run and interpret the outputs of both tools;
- to build a more complete initial subdomain list (**subs.txt**);
- to be able to compare each tool's contributions and identify which one found the most useful domains (discussed later in the Conclusion).

1.2 Methodology

I ran each tool independently, applied lightweight filtering to reduce obvious noise, collected the raw outputs for reproducibility, and then merged and deduplicated the results. I also ran a small statistics script to compare the per-tool results before merging so I could quantify overlap and unique findings.

1.3 Commands

```
1 # Run subfinder
2
3
4 subfinder -d chime.com -silent | sort -u | grep -v '*' >
5     subfinder_out.txt
6
7 # Run assetfinder
8
9 assetfinder --subs-only chime.com | sort -u | grep -v '*' >
10    assetfinder_out.txt
11
12 # Run a small stats script to compare outputs
13
14 python3 stats.py assetfinder_out.txt subfinder_out.txt
15
16 # Merge and deduplicate into passive_subs.txt
17
18 cat subfinder_out.txt assetfinder_out.txt | sort -u > passive_subs.
19     txt
```

1.4 Results

The statistics script produced the following output:

```
1 =====
2 |          STATS          |
3 =====
4
5 - Number of subdomains found
6     assetfinder_out.txt : 85
7     subfinder_out.txt : 269
8
9 - Number of subdomains found
10    By both tools : 77
11    By just assetfinder_out.txt : 8
12    By just subfinder_out.txt : 192
13
14 - Number of INVALID subdomains found (no <chime.com> suffix)
15    By both tools : 0
16    By just assetfinder_out.txt : 0
17    By just subfinder_out.txt : 0
18
19 =====
```

Key observations from the results:

- **subfinder** returned substantially more candidate subdomains than **assetfinder**. Many of **assetfinder**'s results were also present in **subfinder**.
- There were no obviously invalid entries (for example, external domains incorrectly listed as `chime.com` subdomains).
- The overlap (77 entries) indicates a solid common core; the larger number of unique results from **subfinder** warrants manual inspection for quality (quantity does not imply usefulness).

Note: During an earlier run, **assetfinder** returned an unexpected entry (`google.com`), which motivated the validation step.

2 Active subdomain enumeration

2.1 Tools and wordlists

For the active enumeration I used a single tool: **dnsx**. The reason for choosing one tool was pragmatic: most active enumeration tools perform similar DNS probing (the main differences are performance and feature set), so I selected **dnsx** for its simplicity and resolver options. In active enumeration, the choice of wordlist typically has a greater impact on results than the specific tool; for wordlists I used entries from SecLists.

I document below the wordlists used, rationale, and performance considerations:

- SecLists/Discovery/DNS/subdomains-top1million-110000.txt — large, broad list used against the base domain to maximise discovery.
- SecLists/Discovery/DNS/subdomains-top1million-5000.txt — smaller list used when enumerating under each discovered subdomain to limit total query volume.
- Performance notes: applying large wordlists across many parent hosts produces a combinatorial explosion of DNS queries;

2.2 Commands

```
1 # Active enumeration against the base domain
2
3 dnsx -d chime.com -silent -w /SecLists/Discovery/DNS/subdomains-
4     top1million-110000.txt > dnsx_out.txt
5
6 # Active enumeration against each discovered subdomain (brute-
7     forcing children) (high query volume)
8
9 dnsx -d passive_subs.txt -silent -w SecLists/Discovery/DNS/
10    subdomains-top1million-5000.txt > dnsx_out1.txt
11
12 # Merge and deduplicate into passive_subs.txt
13
14 cat dnsx_out.txt dnsx_out1.txt | sort -u > active_subs.txt
```

Notes on the commands above:

- The `-silent` flag reduces noisy output and the results were appended to output files for later analysis.
- Running the 110k-wordlist against the base domain produced many DNS requests (on the order of 110,000 queries) and took a noticeable amount of time; running the 5k-wordlist across the 277 subdomains found earlier resulted in roughly 1.385 million queries and took several hours.
- When doing large active scans, it seems better to: use a fast, reliable resolver pool (`resolvers.txt`), implement rate limiting, and document your scan window and permissions.

2.3 Results and analysis

I compared the `dnsx` outputs with the passive results using the same `stats.py` script used previously.

Example command invocations:

```
1 python3 stats.py dnsx_out.txt subs.txt
2 python3 stats.py dnsx_out1.txt subs.txt
3 python3 stats.py dnsx_out.txt dnsx_out1.txt
```

Summary of findings (condensed):

- `dnsx_out.txt` (base-domain run) returned 32 subdomains, of which 13 were new (not present in `subs.txt`).
- `dnsx_out1.txt` (children-of-subdomains run) returned 26 subdomains, of which 10 were new.
- Intersection between the two `dnsx` runs:

Found by both runs: nautilus.chime.com, dds.chime.com, kyc.chime.com, wireless.chime.com, tags.chime.com

Only in the base-domain run: persona.chime.com, cde.chime.com, mcd.chime.com, notifications.chime.com, beacon.chime.com, lending.chime.com, atomic.chime.com, payday.chime.com

Only in the children-of-subdomains run: `www.workdayxchime.com`, `compass.chime.com`, `access.chime.com`, `bounce.updates.chime.com`, `bounce.accounts.chime.com`

Observations:

- The base-domain brute force found slightly more unique subdomains overall, but probing subdomains-for-children also yielded unique entries that the base run missed. Each approach can be complementary.
- The additional unique hosts discovered by the children-of-subdomains run included some suspicious-looking names (for example, `www.workdayxchime.com`), which may be misconfigurations, CNAMEs, or external services. These should be validated (resolve, check TLS certificate CN/SANs, and inspect HTTP responses) before concluding they are in-scope assets.
- Given the high query count and limited additional value gained from the children-of-subdomains pass in this case, I concluded that it is better deprioritising exhaustive children enumeration unless the base-domain results have been exhausted and more depth is clearly needed.

3 HTTP(S) service discovery

3.1 Tools

For this exercise I only used `httpx`. Before using it however I searched for out of scope domains on bugcrowd (<https://bugcrowd.com/engagements/chime>), copied them into a file called `out_of_scope.txt`, and made sure I didn't target them.

3.2 Commands (examples)

```
1 # Merge and deduplicate passive_subs.txt and active_subs into subs.txt
2
3 cat passive_subs.txt active_subs.txt | sort -u > subs.txt
4
5 # Remove out-of-scope (I created a small python script for that)
6 # The output valid domains are written to the file "subs_in_scope.txt"
7
8 python3 remove_out_of_scope.py subs.txt out_of_scope.txt
9
10 # Probe with httpx (outputs urls.txt)
11
12 httpx -silent -l subs_in_scope.txt -no-fallback -status-code > urls.txt
```

4 URL scanning (detailed analysis)

For this exercise I had to choose one url from the file `urls.txt` created on the previous section. To do so, I followed these steps:

1. Eliminate non-useful responses

- **301 / 307 / 308 (redirects):** Usually just point to canonical domains (e.g. www.chime.com).
 - **403 Forbidden:** Can be useful for fuzzing, but less ideal for initial analysis.
 - **400 Bad Request:** Often unhelpful (may indicate incomplete config or wrong host header).
 - **404 Not Found:** Possible for hidden content, but not good for first choice.
2. Focus on live (200 OK) responses
These are most promising for service and tech fingerprinting:
- https://flashpaper-dev3.chime.com [200]
 - https://docs-dev3.chime.com [200]
 - https://flashpaper.chime.com [200]
 - https://flashpaper-dev1.chime.com [200]
 - https://chime-com-dev1.chime.com [200]
 - https://flashpaper-qa.chime.com [200]
 - https://flashpaper-dev4.chime.com [200]
 - https://vpn.chime.com [200]
 - https://wp-dev2.chime.com [200]
3. Potencial analysis
- **Internal/Dev/QA targets** are often misconfigured and leak more than production.
 - **Docs sites** sometimes expose APIs, Swagger, or internal endpoints.
 - **VPN portals** can leak version info, auth endpoints, or be vulnerable to auth bypass.
 - **Flashpaper** sounds like an internal tool/product, possibly admin dashboard. High chance of findings.
 - **WordPress dev** (wp-dev2). WordPress is notorious for vulnerabilities and may have leftover plugins.

4. Final choice

I choose the url https://wp-dev2.chime.com. The reasons are:

- Dev environment (likely weaker security).
- Built with WordPress (high probability for vunerabilities).
- Good chance of endpoints, debug info, or exposed APIs.

4.1 Top 5 Identified Technologies

For this I used both the Wappalyzer browser extension and nuclei CLI tool. I started by using Wappalyzer, and, although it showed me a lot of technologies being used on the website, it also didn't give me a lot of information about the version of each one. So I tried also using nuclei and run the command:

- nuclei -u https://wp-dev2.chime.com -tags tech

The scan found multiple technologies, mostly WordPress plugins plus infrastructure. After analysing the results, I would say the top 5 findings are the following:

1. WordPress (CMS)

- Detected by: **nuclei & Wappalyzer**

- Multiple plugins detected (Classic Editor, SEO, Super Cache, etc.), some with outdated versions.
- Core CMS platform used to serve the website.

2. PHP (Language)

- Detected by: **Wappalyzer**
- Backend language powering WordPress.
- No version information available.

3. MySQL (Database)

- Detected by: **Wappalyzer**
- Stores site content, user data, and configuration.
- No version information available.

4. Cloudflare (CDN/WAF)

- Detected by: **nuclei & Wappalyzer**
- Provides caching, content delivery network, and web application firewall.
- Masks the origin server IP.

5. Matrix (Messaging Protocol)

- Detected by: **nuclei**
- Integration with external Matrix homeserver (`chime.ems.host`).
- Used for decentralized real-time messaging and federation.

4.2 Content discovery

For this exercise I used ffuf and the wordlist common.txt from SecLists:

Unfortunately, while trying to use ffuf, curl and httpx, it seems I was blocked by their cloudflare security. As such, I continued using a different subdomain that is also interesting, that being <https://flashpaper-dev3.chime.com>.

```

1 # Get error size to filter on ffuf command
2 curl -s -o /dev/null -w "%{http_code} %{size_download}\n" https://
   flashpaper-dev3.chime.com/nonexistent-abc123
3
4 # Quick, low-noise probe (Need to replace <404SIZE> with the value
   got on the first command, in my case was 19)
5 # Changed to only use 5 threads to not alert cloudflare
6 ffuf -u https://flashpaper-dev3.chime.com/FUZZ \
   -w ~/SecLists/Discovery/Web-Content/common.txt \
   -e .php,.html,.htm,.txt,.js \
   -t 5 \
   -mc 200,301,302,403 \
   -fs <404SIZE> \
   -o ffuf-quick.json -of json
7
8
9
10
11
12
13

```

The scan yielded multiple results, mostly with http status code 403. From the results, there were only found 4 endpoints with status code 200, those being:

- <https://flashpaper-dev3.chime.com/add>

- <https://flashpaper-dev3.chime.com/favicon.ico>
- <https://flashpaper-dev3.chime.com/ping>
- <https://flashpaper-dev3.chime.com/.well-known/http-opportunistic>

None of them being particularly useful...